

```
pip install datasets
```

```
Requirement already satisfied: datasets in /usr/local/lib/python3.10/dist-packages (3.1.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from datasets) (3.16.1)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from datasets) (1.26.4)
Requirement already satisfied: pyarrow>=15.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (17.0.0)
Requirement already satisfied: dill<0.3.9,>=0.3.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.3.8)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from datasets) (2.2.2)
Requirement already satisfied: requests>=2.32.2 in /usr/local/lib/python3.10/dist-packages (from datasets) (2.32.3)
Requirement already satisfied: tqdm>=4.66.3 in /usr/local/lib/python3.10/dist-packages (from datasets) (4.66.6)
Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages (from datasets) (3.5.0)
Requirement already satisfied: multiprocess<0.70.17 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.70.16)
Requirement already satisfied: fsspec<=2024.9.0,>=2023.1.0 in /usr/local/lib/python3.10/dist-packages (from fsspec[http]<=2024.9.0,>=2023.1.0) (2024.9.0)
Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets) (3.11.2)
Requirement already satisfied: huggingface-hub>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.26.2)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from datasets) (24.2)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from datasets) (6.0.2)
Requirement already satisfied: aiohappyeyeballs>=2.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (2.4.3)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.3.1)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (24.2.0)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.5.0)
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (6.1.0)
Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (0.2.0)
Requirement already satisfied: yarl<2.0,>=1.17.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.17.2)
Requirement already satisfied: async-timeout<6.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (4.0.3)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.23.0->dataclasses-json) (4.12.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.32.2->datasets) (3.3.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.32.2->datasets) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.32.2->datasets) (2.2.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.32.2->datasets) (2024.8.30)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2024.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->datasets) (1.16.0)
```

```
import pandas as pd
from sklearn.model_selection import train_test_split
from datasets import Dataset

# Load dataset
data = pd.read_csv("bbc_text_cls.csv") # Replace with the actual dataset path
data = data[['text', 'labels']] # Ensure "labels" column name is accurate
```

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Define the label-to-category mapping
label_mapping = {
    0: "Business",
    1: "Entertainment",
    2: "Politics",
    3: "Sport",
    4: "Tech"
}

# Assuming `data` is your DataFrame containing the numeric labels in 'label' column
# Map numeric labels to category names
data['category'] = data['labels'].replace(label_mapping)

# Calculate the count of each category
label_counts = data['category'].value_counts()


# Enhanced bar chart using Matplotlib
plt.figure(figsize=(10, 6))
bars = plt.bar(label_counts.index, label_counts.values, color=sns.color_palette("viridis", len(label_counts)))

# Add value annotations above each bar
for bar in bars:
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, yval + 1, int(yval), ha='center', va='bottom', fontsize=12)

# Customize the chart
plt.xlabel("Categories", fontsize=14)
```

```
plt.ylabel("Count", fontsize=14)
plt.title("Distribution of Categories", fontsize=16, fontweight="bold")
plt.xticks(rotation=0, fontsize=12)
plt.yticks(fontsize=12)

plt.tight_layout()
plt.show()
```



```
# Encode labels to numerical values
data['label'] = data['labels'].astype('category').cat.codes

# Step 1: Split into training and test sets (80-20 split) - test set is set aside for final evaluation
train_data, test_data = train_test_split(data, test_size=0.2, random_state=42, stratify=data['label'])

# Step 2: Further split train_data into training and validation sets (80-20 split of training data)
train_texts, val_texts, train_labels, val_labels = train_test_split(
    train_data['text'], train_data['label'], test_size=0.2, random_state=42, stratify=train_data['label']
)

# Step 3: Convert to Hugging Face Dataset format
train_data = Dataset.from_dict({'text': train_texts, 'label': train_labels})
val_data = Dataset.from_dict({'text': val_texts, 'label': val_labels})
test_data = Dataset.from_dict({'text': test_data['text'], 'label': test_data['label']})

from transformers import DistilBertTokenizer

# Load DistilBERT tokenizer
tokenizer = DistilBertTokenizer.from_pretrained('distilbert-base-uncased')

# Tokenize function
def tokenize_function(example):
    return tokenizer(example['text'], padding='max_length', truncation=True, max_length=128)

# Tokenize datasets
train_data = train_data.map(tokenize_function, batched=True)
val_data = val_data.map(tokenize_function, batched=True)
test_data = test_data.map(tokenize_function, batched=True)

# Set the format for PyTorch
train_data.set_format(type='torch', columns=['input_ids', 'attention_mask', 'label'])
val_data.set_format(type='torch', columns=['input_ids', 'attention_mask', 'label'])
test_data.set_format(type='torch', columns=['input_ids', 'attention_mask', 'label'])
```

Map: 100% 1424/1424 [00:12<00:00, 116.42 examples/s]

Map: 100% 356/356 [00:03<00:00, 114.53 examples/s]

Map: 100% 145/145 [00:03<00:00, 116.51 examples/s]

```
from transformers import DistilBertForSequenceClassification
```

```
# Load DistilBERT model for classification with 5 labels
```

```
model = DistilBertForSequenceClassification.from_pretrained('distilbert-base-uncased', num_labels=5)
```

```
# Freeze all DistilBERT layers except the classification head
```

```
for param in model.distilbert.parameters():
    param.requires_grad = False
```

Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased and are new. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
# Define metrics function with scalar conversion
```

```
from sklearn.metrics import precision_recall_fscore_support
```

```
def compute_metrics(pred):
```

```
    labels = pred.label_ids
```

```
    preds = pred.predictions.argmax(-1)
```

```
    # Calculate per-class metrics
```

```
    precision, recall, f1, _ = precision_recall_fscore_support(labels, preds, average=None)
```

```
    # Print per-class metrics
```

```
    for i in range(len(precision)):
```

```
        print(f"Class {i} - Precision: {precision[i]:.4f}, Recall: {recall[i]:.4f}, F1: {f1[i]:.4f}")
```

```
    # Calculate average metrics for logging
```

```
    avg_precision = sum(precision) / len(precision)
```

```
    avg_recall = sum(recall) / len(recall)
```

```
    avg_f1 = sum(f1) / len(f1)
```

```
    # Return average metrics for logging
```

```
    return {
```

```
        'precision': avg_precision,
```

```
        'recall': avg_recall,
```

```
        'f1': avg_f1
```

```
    }
```

```
# Set training arguments
```

```
from transformers import TrainingArguments
```

```
training_args = TrainingArguments(
```

```
    output_dir='./results',
```

```
    evaluation_strategy="epoch",          # Evaluate model every epoch
```

```
    save_strategy="epoch",               # Save model every epoch
```

```
    logging_strategy="epoch",           # Log metrics every epoch
```

```
    per_device_train_batch_size=16,     # Adjust batch size as needed
```

```
    per_device_eval_batch_size=64,      # Evaluation batch size
```

```
    num_train_epochs=30,                # Maximum number of epochs
```

```
    save_total_limit=2,                 # Limit total saved models
```

```
    load_best_model_at_end=True,        # Load the best model at the end of training
```

```
    metric_for_best_model="eval_f1",    # Use F1 as the evaluation metric
```

```
    greater_is_better=True,            # Higher F1 is better
```

```
    logging_dir='./logs',
```

```
    learning_rate=5e-5,
```

```
    report_to="none"                   # Suppress default logging to external platforms
```

```
)
```

/usr/local/lib/python3.10/dist-packages/transformers/training_args.py:1568: FutureWarning: `evaluation_strategy` is deprecated and will be removed in a future version. Use `eval_strategy` instead.

```
# Initialize and train the Trainer
from transformers import EarlyStoppingCallback, Trainer, TrainingArguments
from transformers import Trainer
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_data,
    eval_dataset=val_data,
    compute_metrics=compute_metrics,
    callbacks=[EarlyStoppingCallback(early_stopping_patience=3)])
```

```
# Train the model
```

```
trainer.train()
```

```
# Evaluate the model
```

```
test_results = trainer.evaluate(test_data)
```

```
print("Test Set Evaluation:", test_results)
```

 [534/2670 00:15 < 01:01, 34.78 it/s, Epoch 6/30]

Epoch	Training Loss	Validation Loss	Precision	Recall	F1
1	1.390600	1.142759	0.906817	0.849452	0.858851
2	0.944400	0.739522	0.944593	0.924695	0.930341
3	0.603300	0.473422	0.958824	0.956637	0.957588
4	0.405100	0.339111	0.948384	0.951479	0.949549
5	0.293400	0.258226	0.956863	0.954299	0.955212
6	0.236400	0.217510	0.949358	0.950107	0.949596

Class 0 - Precision: 0.6864, Recall: 0.9878, F1: 0.8100

Class 1 - Precision: 0.9787, Recall: 0.7419, F1: 0.8440

Class 2 - Precision: 0.9841, Recall: 0.9394, F1: 0.9612

Class 3 - Precision: 0.9111, Recall: 1.0000, F1: 0.9535

Class 4 - Precision: 0.9737, Recall: 0.5781, F1: 0.7255

Class 0 - Precision: 0.7980, Recall: 0.9634, F1: 0.8729

Class 1 - Precision: 0.9833, Recall: 0.9516, F1: 0.9672

Class 2 - Precision: 0.9538, Recall: 0.9394, F1: 0.9466

Class 3 - Precision: 0.9878, Recall: 0.9878, F1: 0.9878

Class 4 - Precision: 1.0000, Recall: 0.7812, F1: 0.8772

Class 0 - Precision: 0.9167, Recall: 0.9390, F1: 0.9277

Class 1 - Precision: 0.9677, Recall: 0.9677, F1: 0.9677

Class 2 - Precision: 0.9545, Recall: 0.9545, F1: 0.9545

Class 3 - Precision: 0.9880, Recall: 1.0000, F1: 0.9939

Class 4 - Precision: 0.9672, Recall: 0.9219, F1: 0.9440

Class 0 - Precision: 0.9342, Recall: 0.8659, F1: 0.8987

Class 1 - Precision: 0.9683, Recall: 0.9839, F1: 0.9760

Class 2 - Precision: 0.9130, Recall: 0.9545, F1: 0.9333

Class 3 - Precision: 0.9880, Recall: 1.0000, F1: 0.9939

Class 4 - Precision: 0.9385, Recall: 0.9531, F1: 0.9457

Class 0 - Precision: 0.9048, Recall: 0.9268, F1: 0.9157

Class 1 - Precision: 0.9683, Recall: 0.9839, F1: 0.9760

Class 2 - Precision: 0.9403, Recall: 0.9545, F1: 0.9474

Class 3 - Precision: 0.9880, Recall: 1.0000, F1: 0.9939

Class 4 - Precision: 0.9831, Recall: 0.9062, F1: 0.9431

Class 0 - Precision: 0.9125, Recall: 0.8902, F1: 0.9012

Class 1 - Precision: 0.9683, Recall: 0.9839, F1: 0.9760

Class 2 - Precision: 0.9265, Recall: 0.9545, F1: 0.9403

Class 3 - Precision: 0.9880, Recall: 1.0000, F1: 0.9939

Class 4 - Precision: 0.9516, Recall: 0.9219, F1: 0.9365

[7/7 00:00]

Class 0 - Precision: 0.9412, Recall: 0.9412, F1: 0.9412

Class 1 - Precision: 0.9868, Recall: 0.9740, F1: 0.9804

Class 2 - Precision: 0.9310, Recall: 0.9643, F1: 0.9474

Class 3 - Precision: 0.9709, Recall: 0.9804, F1: 0.9756

Class 4 - Precision: 0.9870, Recall: 0.9500, F1: 0.9682

Test Set Evaluation: {'eval_loss': 0.4489331543445587, 'eval_precision': 0.9633879664061535, 'eval_recall': 0.9619760631525338, 'eval_f1':

```
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Generate predictions
```

```
predictions = trainer.predict(test_data)
```

```

pred_labels = predictions.predictions.argmax(-1)
true_labels = predictions.label_ids

# Print classification report
class_labels = ['business', 'entertainment', 'politics', 'sport', 'tech']
print("\nClassification Report:\n", classification_report(true_labels, pred_labels, target_names=class_labels))

```

```

↗ Class 0 - Precision: 0.9412, Recall: 0.9412, F1: 0.9412
Class 1 - Precision: 0.9868, Recall: 0.9740, F1: 0.9804
Class 2 - Precision: 0.9310, Recall: 0.9643, F1: 0.9474
Class 3 - Precision: 0.9709, Recall: 0.9804, F1: 0.9756
Class 4 - Precision: 0.9870, Recall: 0.9500, F1: 0.9682

```

```

Classification Report:

```

	precision	recall	f1-score	support
business	0.94	0.94	0.94	102
entertainment	0.99	0.97	0.98	77
politics	0.93	0.96	0.95	84
sport	0.97	0.98	0.98	102
tech	0.99	0.95	0.97	80
accuracy			0.96	445
macro avg	0.96	0.96	0.96	445
weighted avg	0.96	0.96	0.96	445

```


def plot_confusion_matrix(true_labels, pred_labels, classes):
    cm = confusion_matrix(true_labels, pred_labels)
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=classes, yticklabels=classes)
    plt.xlabel('Predicted Label')
    plt.ylabel('True Label')
    plt.title('Confusion Matrix')
    plt.show()

```

```

# Define class names based on your labels
class_names = ['Business', 'Entertainment', 'Politics', 'Sport', 'Tech']
plot_confusion_matrix(true_labels, pred_labels, classes=class_names)

```



Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

```

import matplotlib.pyplot as plt


# Initialize lists to store epoch-wise metrics
epochs = []
train_loss = []
val_loss = []
val_f1 = []
train_f1 = []

# Extract metrics from the trainer's log history
for log in trainer.state.log_history:
    if 'epoch' in log:
        current_epoch = log['epoch']
        if 'loss' in log:
            train_loss.append(log['loss'])
            epochs.append(current_epoch)
        if 'eval_loss' in log:
            val_loss.append(log['eval_loss'])
        if 'eval_f1' in log:
            val_f1.append(log['eval_f1'])

# Align lengths by truncating to the shortest list length
min_len = min(len(epochs), len(train_loss), len(val_loss), len(val_f1))
epochs = epochs[:min_len]
train_loss = train_loss[:min_len]
val_loss = val_loss[:min_len]
val_f1 = val_f1[:min_len]

# Plot Training and Validation Loss over Epochs
plt.figure(figsize=(10, 5))
plt.plot(epochs, train_loss, marker='o', label="Training Loss")
plt.plot(epochs, val_loss, marker='o', color='orange', label="Validation Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Training and Validation Loss over Epochs")
plt.legend()
plt.show()

```



```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load dataset
data = pd.read_csv("bbc_text_cls.csv") # Replace with the correct path to your dataset

# Map numeric labels to category names if needed
label_mapping = {0: "Business", 1: "Entertainment", 2: "Politics", 3: "Sport", 4: "Tech"}
data['category'] = data['labels'].replace(label_mapping)

```


```
# Count categories
label_counts = data['category'].value_counts()

# Plot bar chart
plt.figure(figsize=(10, 6))
sns.barplot(x=label_counts.index, y=label_counts.values, palette="viridis")
plt.title("Distribution of Categories", fontsize=16)
plt.xlabel("Category", fontsize=14)
plt.ylabel("Count", fontsize=14)
plt.xticks(fontsize=12)
plt.tight_layout()
plt.show()
```

↗ <ipython-input-39-1d3599f17167>:17: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

```
sns.barplot(x=label_counts.index, y=label_counts.values, palette="viridis")
```



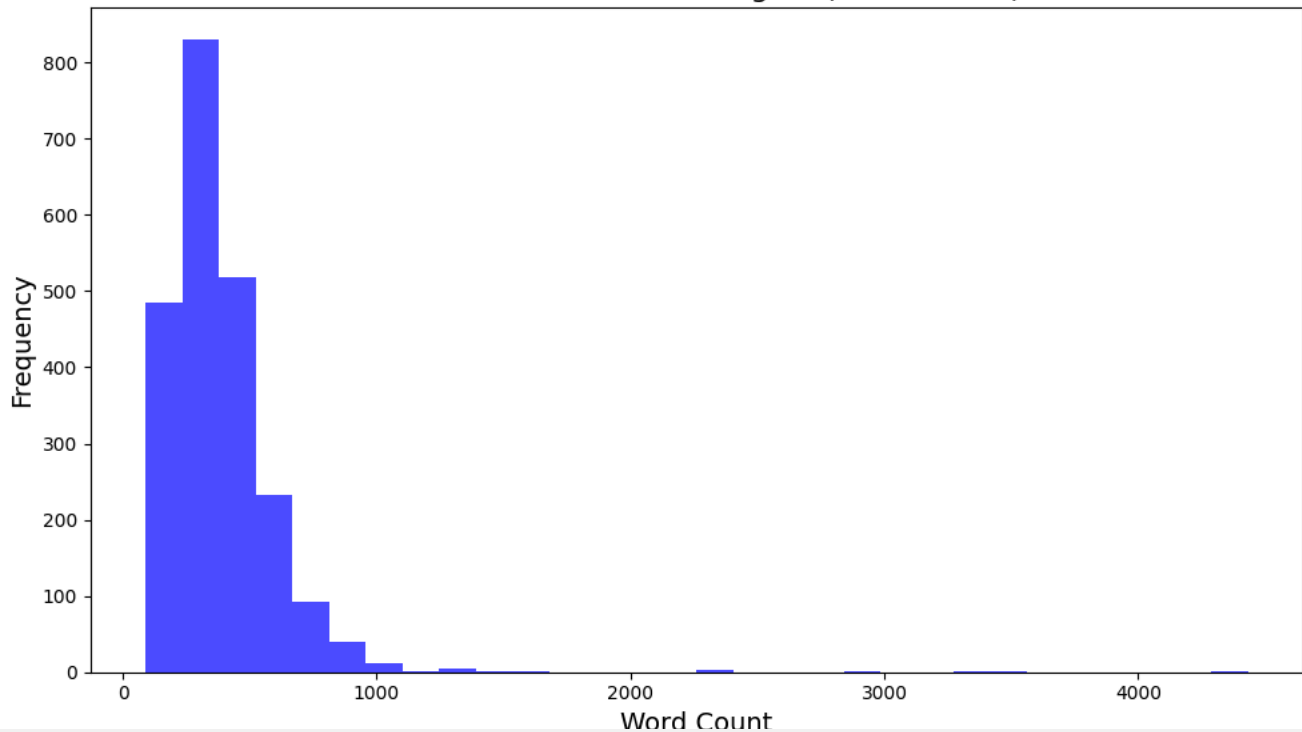
Start coding or [generate](#) with AI.

```
# Calculate text lengths
data['text_length'] = data['text'].apply(lambda x: len(x.split()))

# Plot histogram
plt.figure(figsize=(10, 6))
plt.hist(data['text_length'], bins=30, color='blue', alpha=0.7)
plt.title("Distribution of Article Lengths (Word Count)", fontsize=16)
plt.xlabel("Word Count", fontsize=14)
plt.ylabel("Frequency", fontsize=14)
plt.tight_layout()
plt.show()
```



Distribution of Article Lengths (Word Count)



```
from sklearn.model_selection import train_test_split


# Split the dataset
train_data, test_data = train_test_split(data, test_size=0.2, stratify=data['category'], random_state=42)
train_data, val_data = train_test_split(train_data, test_size=0.2, stratify=train_data['category'], random_state=42)

# Sizes for pie chart
sizes = [len(train_data), len(val_data), len(test_data)]
labels = ["Train", "Validation", "Test"]

# Plot pie chart
plt.figure(figsize=(8, 8))
plt.pie(sizes, labels=labels, autopct='%1.1f%%', startangle=140, colors=['skyblue', 'lightgreen', 'salmon'])
plt.title("Train-Validation-Test Split", fontsize=16)
plt.show()
```




Train-Validation-Test Split




```
# Calculate text lengths
data['text_length'] = data['text'].apply(lambda x: len(x.split()))

# Plot histogram with adjusted range and bins
plt.figure(figsize=(10, 6))
plt.hist(data['text_length'], bins=30, range=(data['text_length'].min(), data['text_length'].max()),
         color='blue', alpha=0.7, edgecolor='black')
plt.title("Distribution of Article Lengths (Word Count)", fontsize=16)
plt.xlabel("Word Count", fontsize=14)
plt.ylabel("Frequency", fontsize=14)
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)
plt.tight_layout()
plt.show()
```



Distribution of Article Lengths (Word Count)



```
data.shape
```

```
(2225, 4)
```

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Load dataset
data = pd.read_csv("bbc_text_cls.csv") # Replace with the correct path to your dataset
```

```
# Map numeric labels to category names if needed
label_mapping = {0: "Business", 1: "Entertainment", 2: "Politics", 3: "Sport", 4: "Tech"}
data['category'] = data['labels'].replace(label_mapping)
```

```
# Count categories
label_counts = data['category'].value_counts()
```

```
# Define medium-tone colors for each category
colors = ["#E57373", "#FBC02D", "#64B5F6", "#81C784", "#BA68C8"] # Softer shades
```

```
# Plot bar chart
plt.figure(figsize=(10, 6))
bars = sns.barplot(x=label_counts.index, y=label_counts.values, palette=colors)
plt.title("Distribution of Categories", fontsize=16)
plt.xlabel("Category", fontsize=14)
plt.ylabel("Count", fontsize=14)
plt.xticks(fontsize=12)
```

```
# Add labels above each bar
for bar in bars.patches:
    bar_height = bar.get_height()
    bars.annotate(f"{int(bar_height)}",
                  (bar.get_x() + bar.get_width() / 2, bar_height),
                  ha='center', va='bottom', fontsize=12, color='black')
```

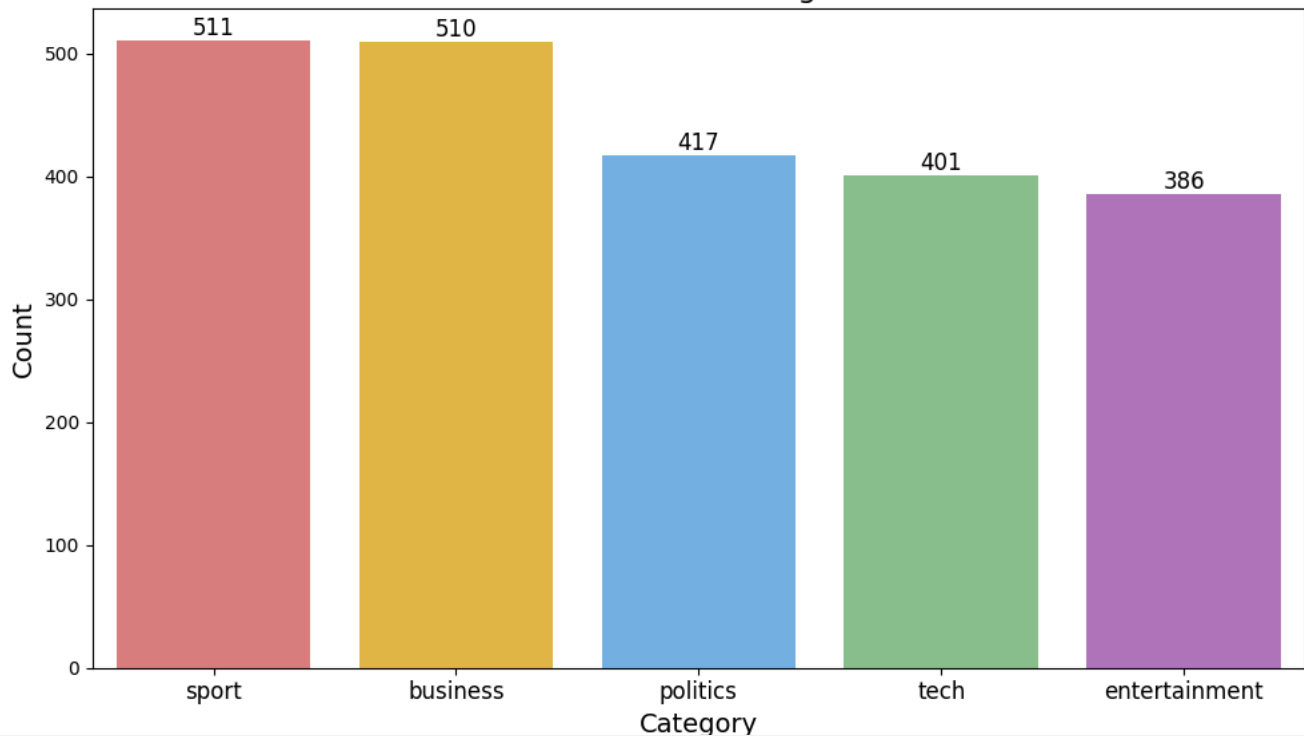
```
plt.tight_layout()
plt.show()
```

```
<ipython-input-44-07e54f3a2887>:20: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

```
bars = sns.barplot(x=label_counts.index, y=label_counts.values, palette=colors)
```

Distribution of Categories



```
# Calculate text lengths
data['text_length'] = data['text'].apply(lambda x: len(x.split()))

# Calculate statistics
mean_length = data['text_length'].mean()
median_length = data['text_length'].median()

# Plot histogram
plt.figure(figsize=(12, 6))
plt.hist(data['text_length'], bins=30, color='skyblue', alpha=0.7, edgecolor='black')
plt.axvline(mean_length, color='red', linestyle='--', label=f'Mean: {mean_length:.2f}')
plt.axvline(median_length, color='green', linestyle='-', label=f'Median: {median_length:.0f}')

# Add gridlines
plt.grid(axis='y', linestyle='--', alpha=0.7)

# Add title and labels
plt.title("Distribution of Article Lengths (Word Count)", fontsize=18, fontweight='bold')
plt.xlabel("Word Count", fontsize=14)
plt.ylabel("Frequency", fontsize=14)


# Add legend
plt.legend(fontsize=12)

# Customize ticks
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)

# Adjust layout
plt.tight_layout()
plt.show()
```



Distribution of Article Lengths (Word Count)



Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

```
# Confirm which parameters are trainable
for name, param in model.named_parameters():
    print(f"{name}: requires_grad = {param.requires_grad}")
```

```
distilbert.embeddings.word_embeddings.weight: requires_grad = False
distilbert.embeddings.position_embeddings.weight: requires_grad = False
distilbert.embeddings.LayerNorm.weight: requires_grad = False
distilbert.embeddings.LayerNorm.bias: requires_grad = False
distilbert.transformer.layer.0.attention.q_lin.weight: requires_grad = False
distilbert.transformer.layer.0.attention.q_lin.bias: requires_grad = False
distilbert.transformer.layer.0.attention.k_lin.weight: requires_grad = False
distilbert.transformer.layer.0.attention.k_lin.bias: requires_grad = False
distilbert.transformer.layer.0.attention.v_lin.weight: requires_grad = False
```